






Gdev: First-Class GPU Resource Management in the Operating System

Shinpei Kato
Nagoya University

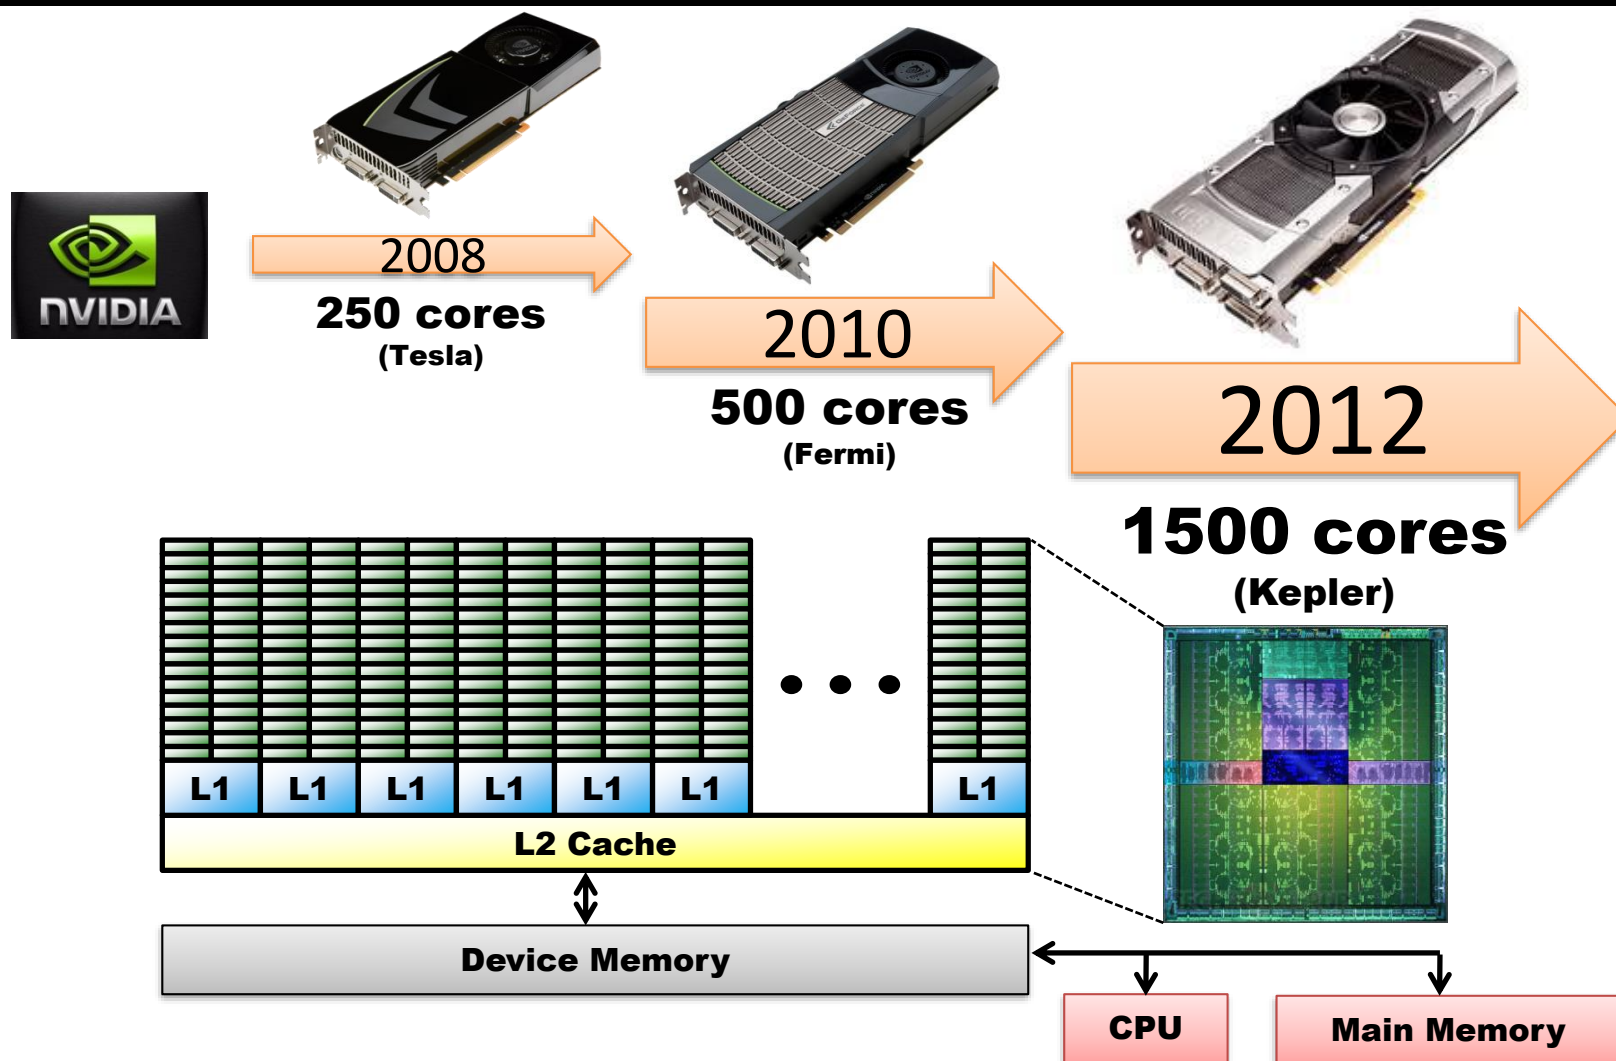
Michael McThrow Carlos Maltzahn Scott Brandt
UC Santa Cruz




GPUs embrace “many cores”.




Graphics Processing Unit (GPU)



Performance Trend




Single Precision Performance




Performance per Watt

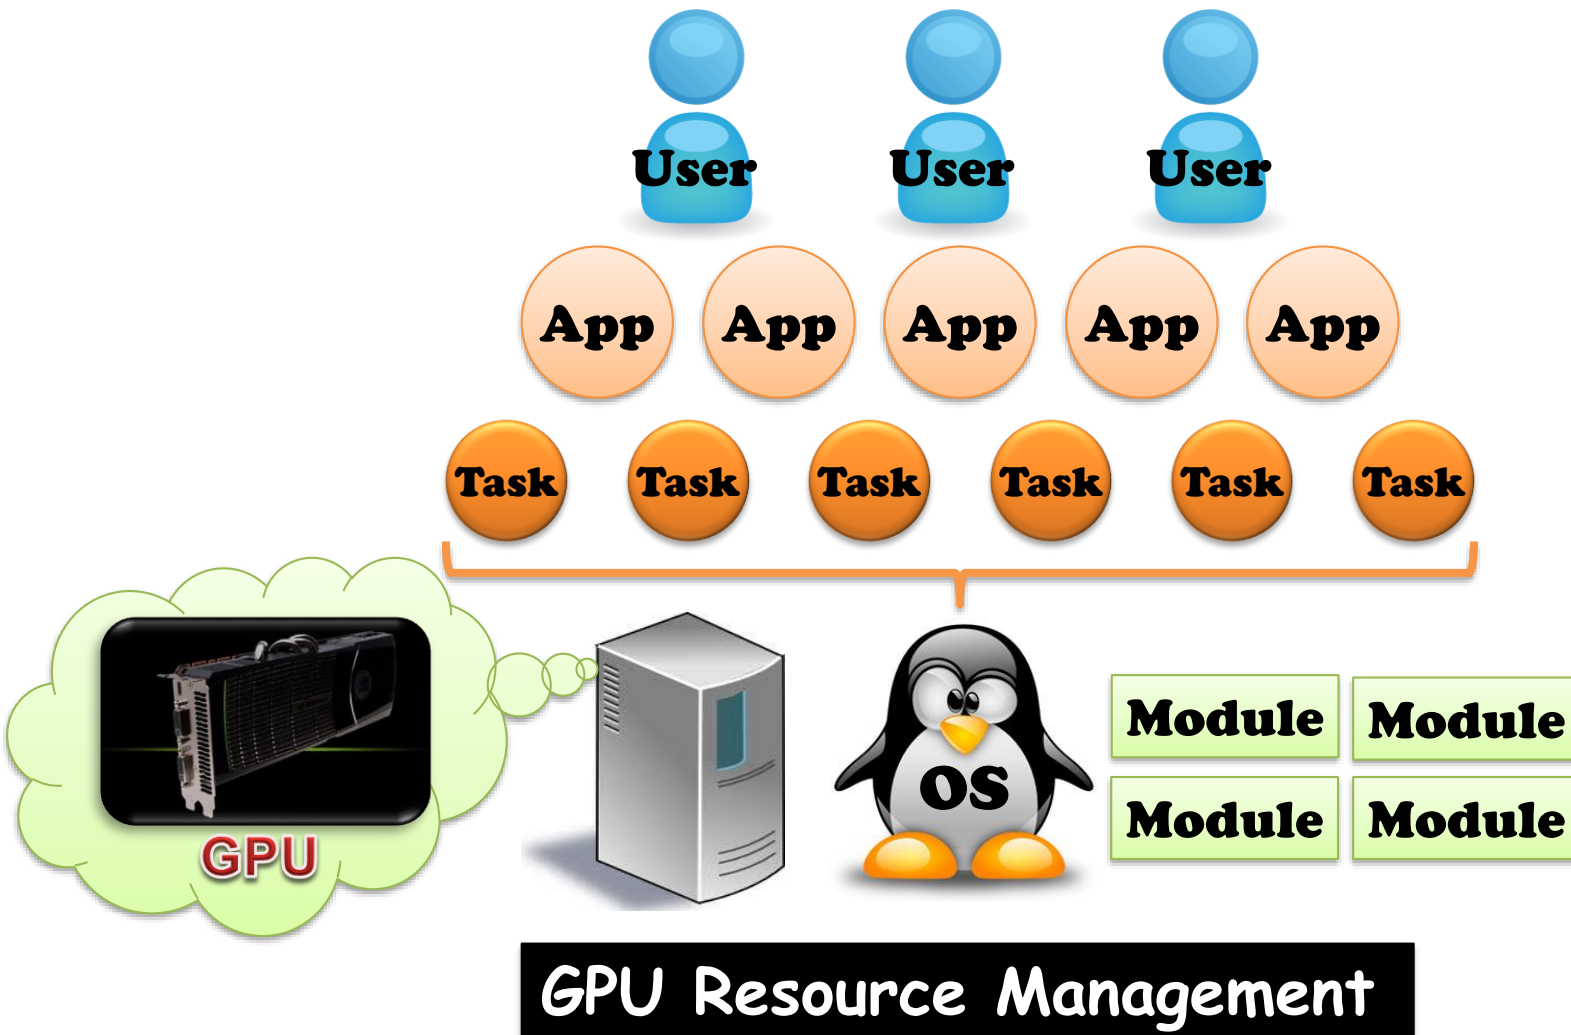
GPUs Suit Science



Traffic Simulation



Not Yet “General-Purpose”



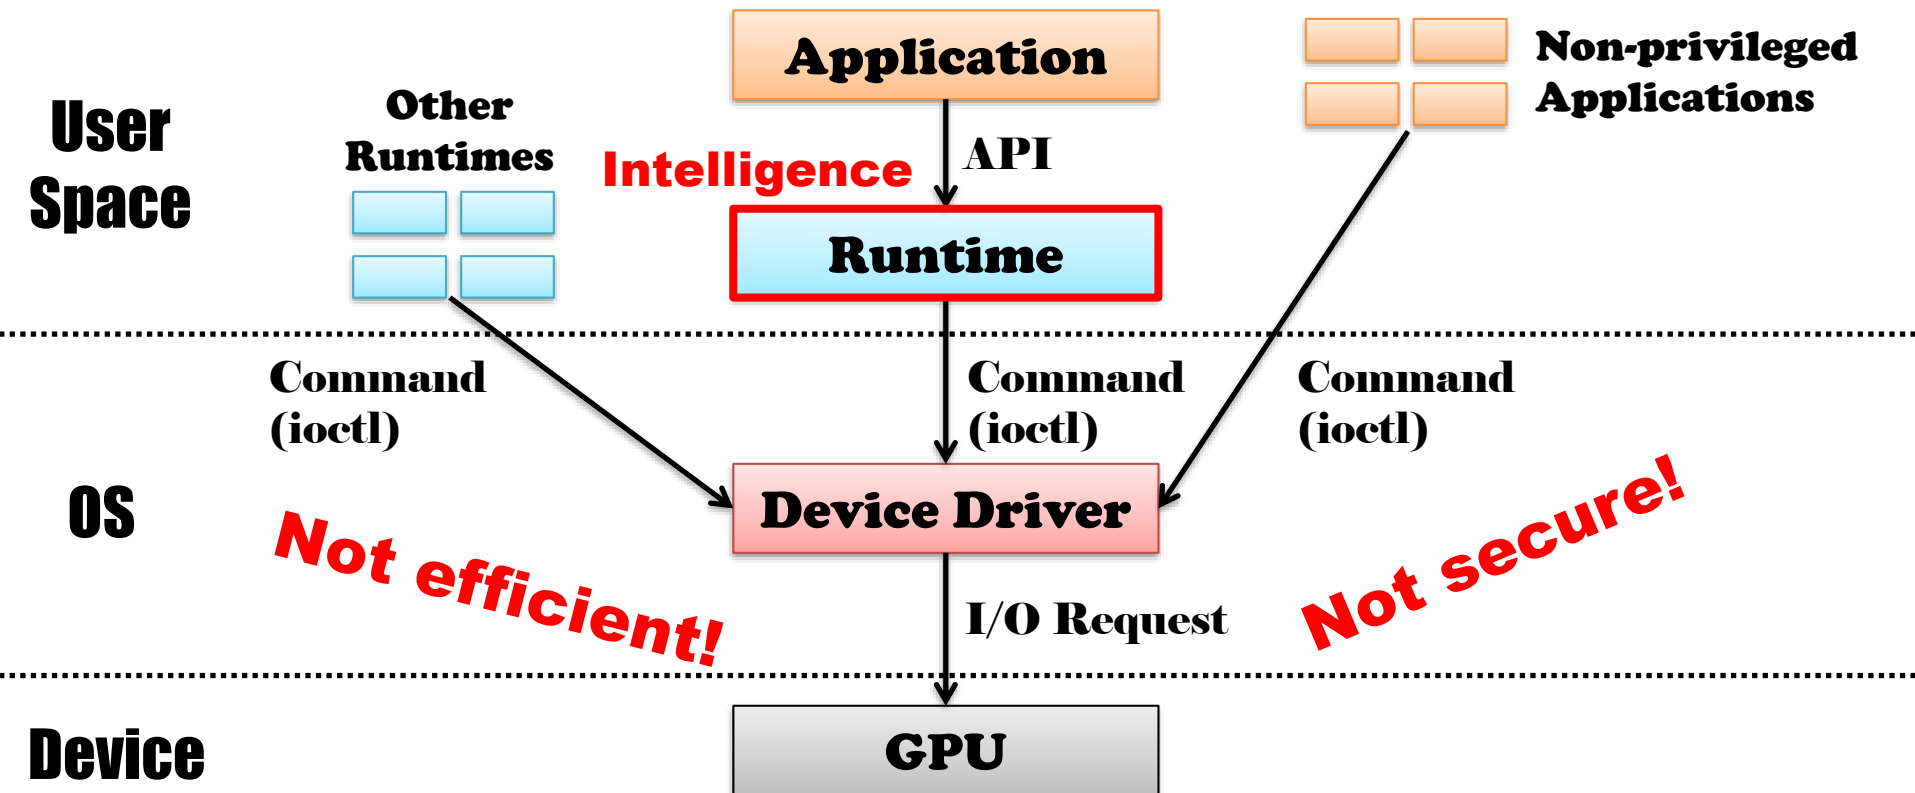
Gdev

- **New approach** to GPU resource management
 - Allows the OS as well as user-space applications to use GPUs.
- **New functions** of GPU resource management
 - Shared device memory (IPC)
 - Data swapping
 - System-level virtualization
- **Open-source implementation**

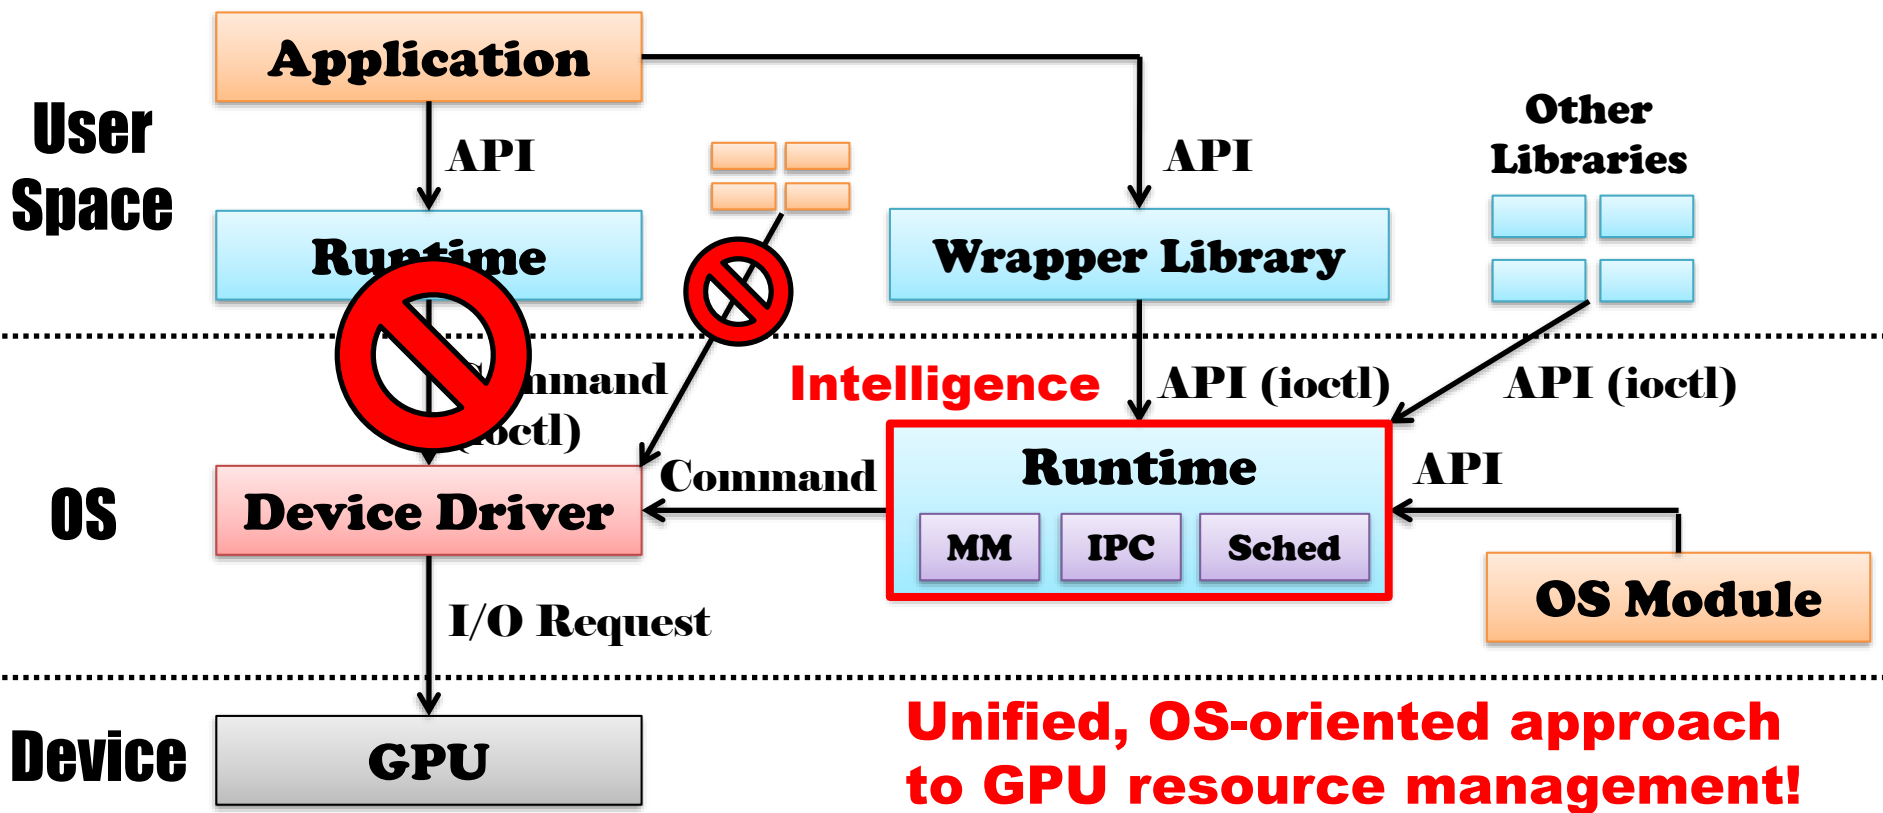
Outline

- **Motivation**
- **Approach**
- **GPU Resource Management**
- **Evaluation**
- **Conclusion**

Traditional Naïve Approach



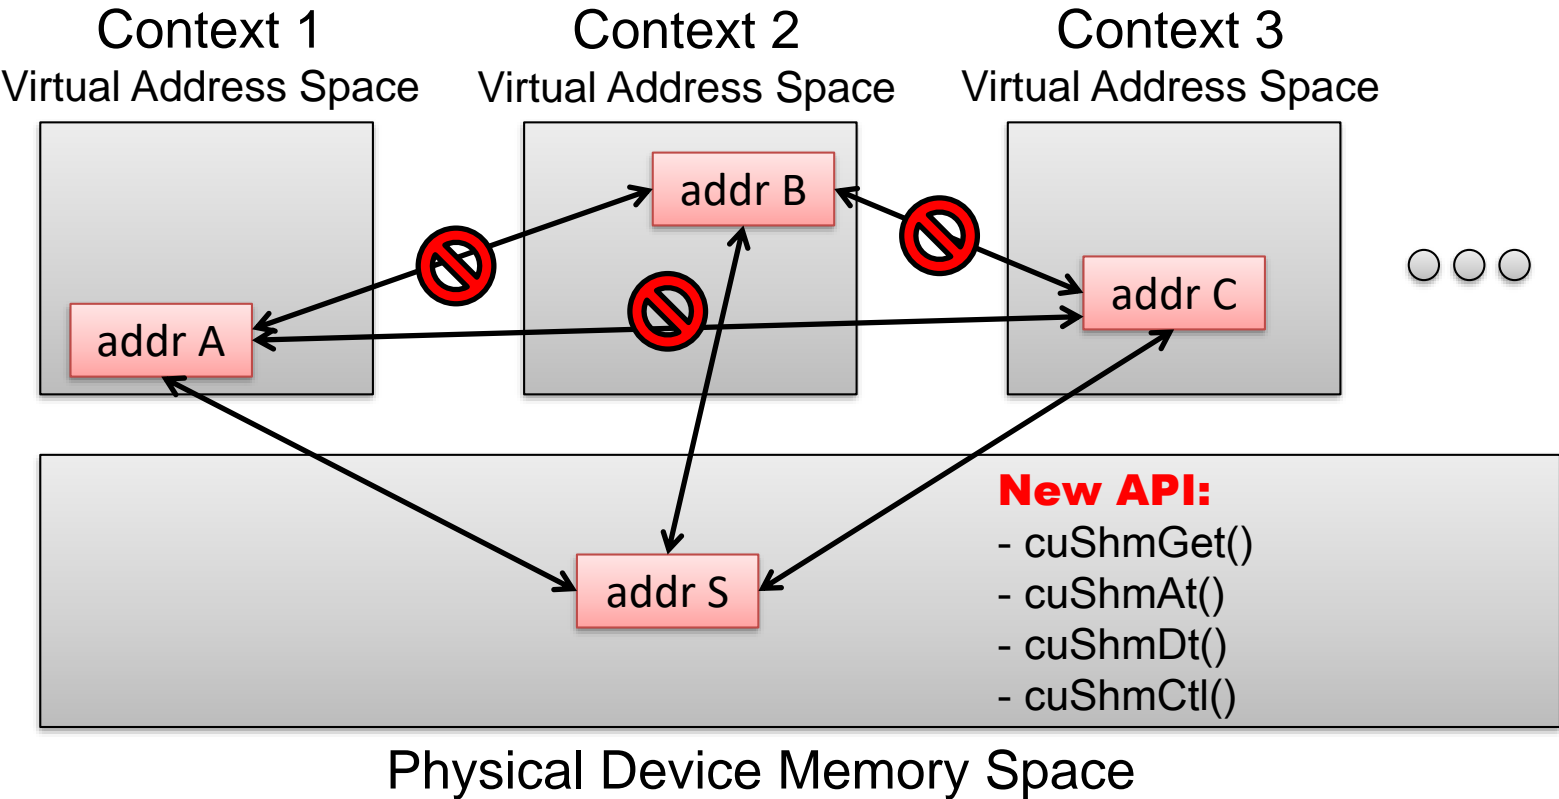
Gdev Approach



Outline



- Motivation
- Approach
- **GPU Resource Management**
- **Evaluation**
- **Conclusion**

Shared Device Memory




E.g., Dataflow (2x2 Tree)


No Shared Memory



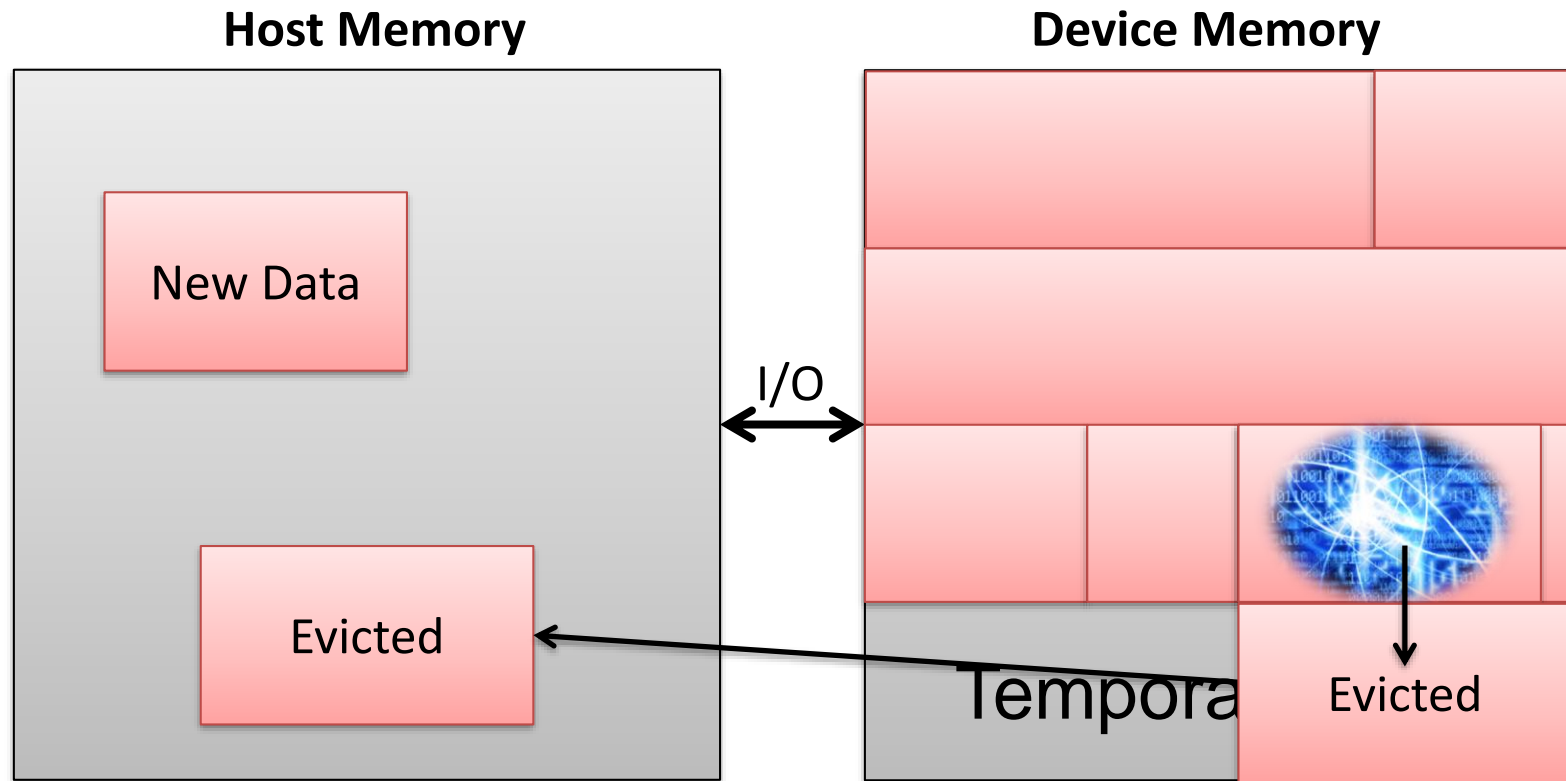
Shared Memory IPC




Data Swapping



Data Swapping (Enhanced)

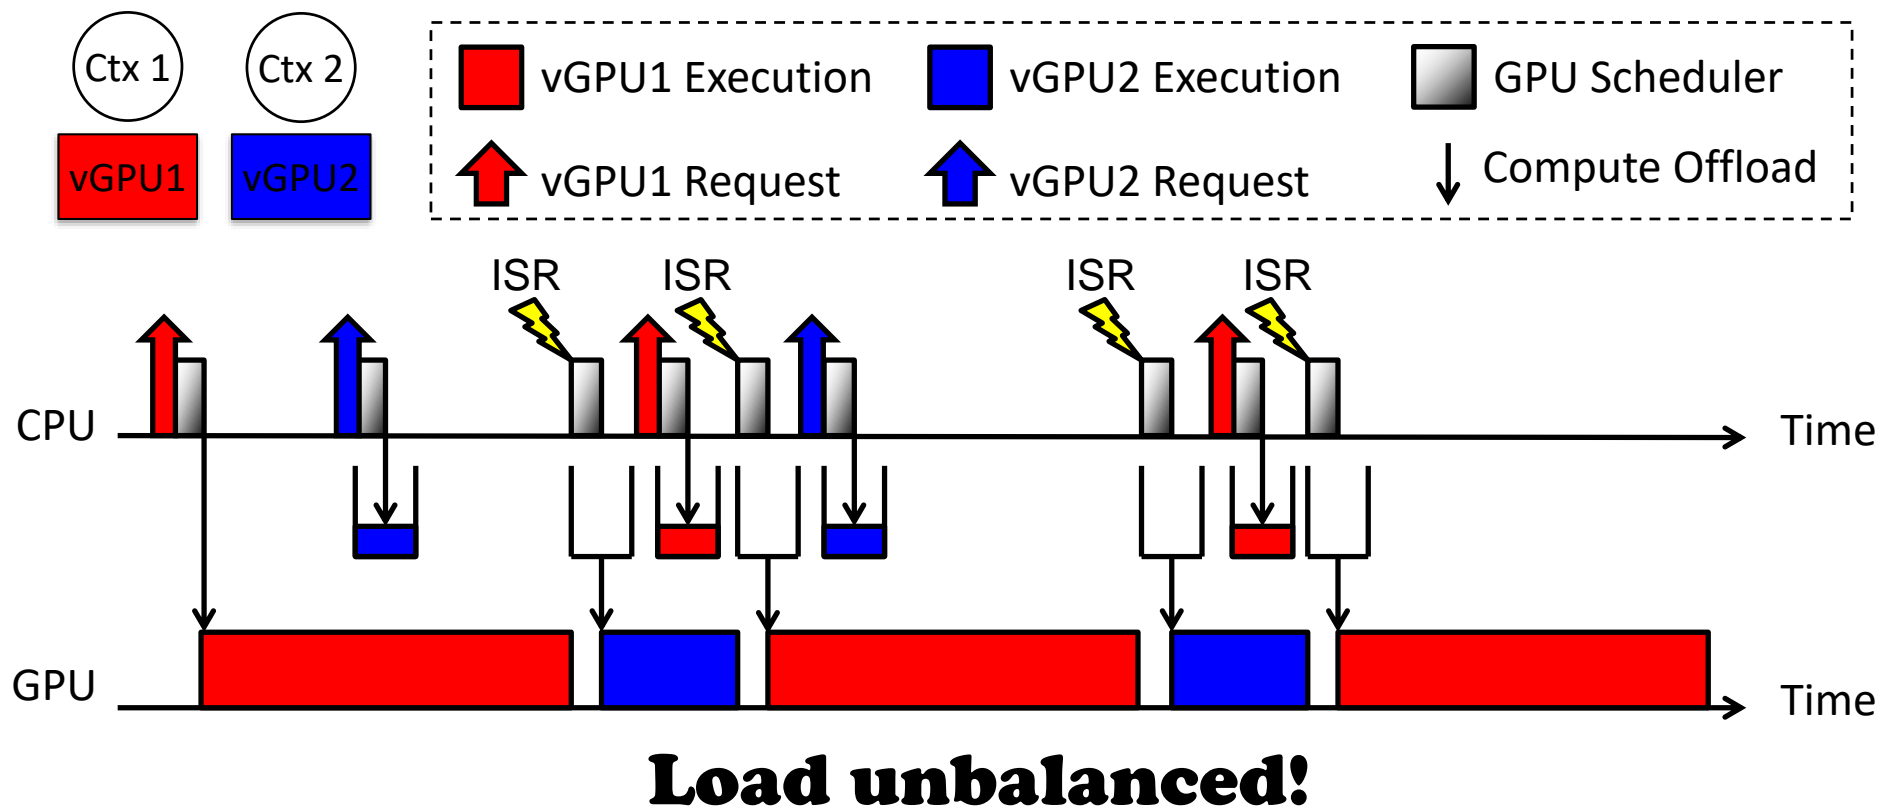


GPU Virtualization

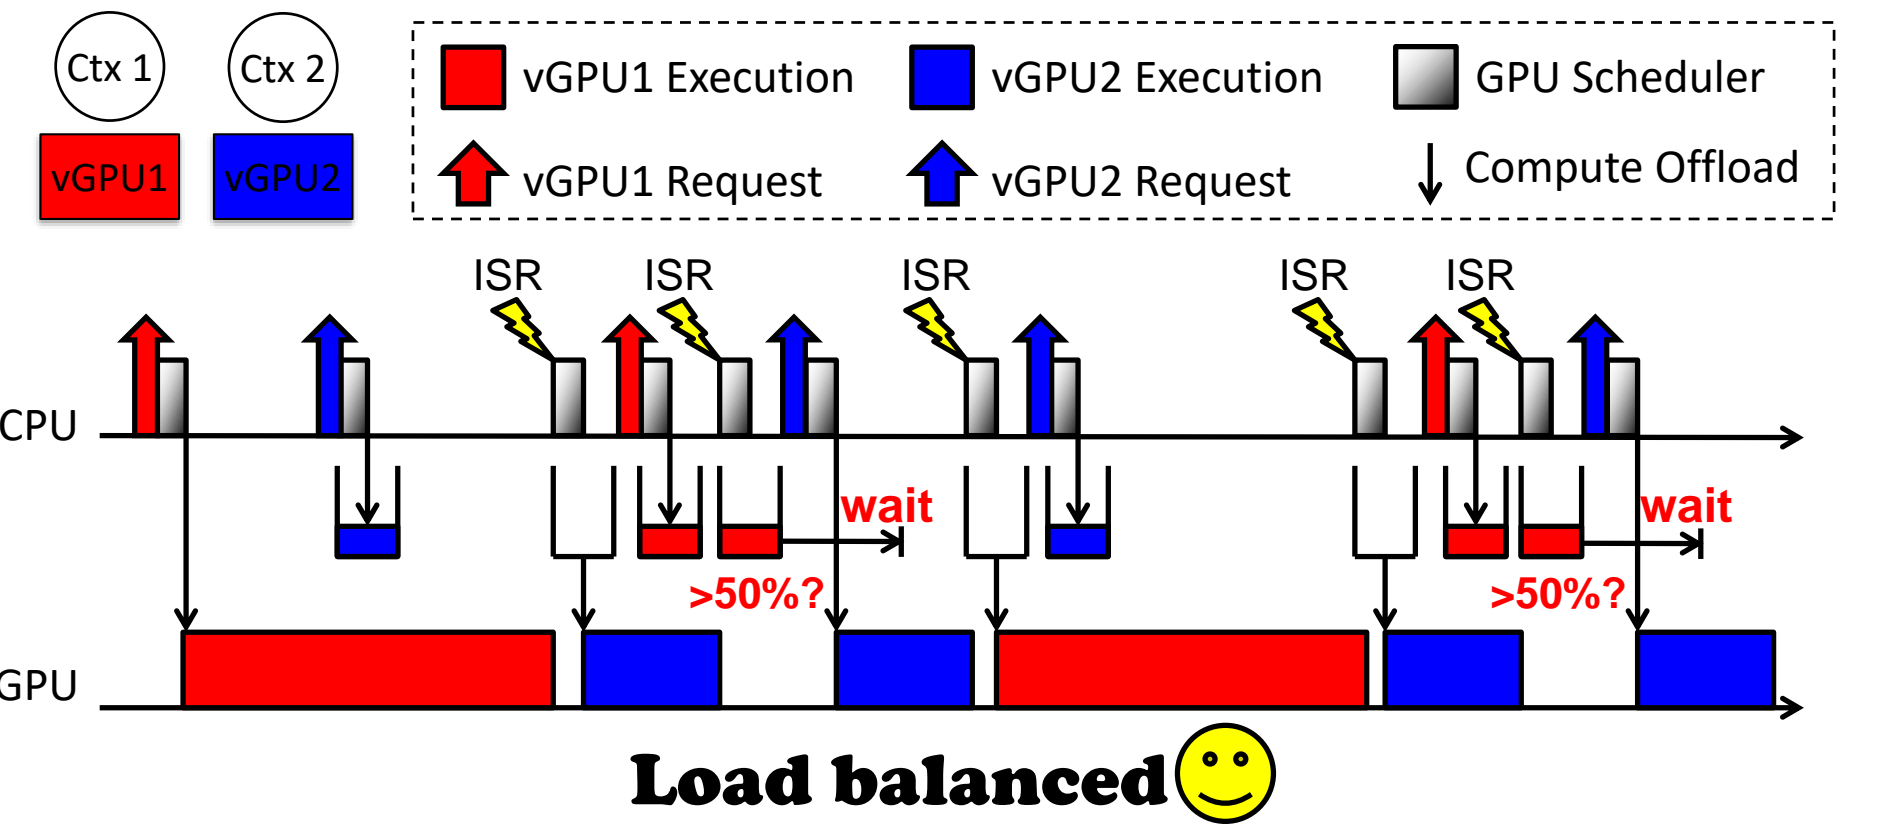


Existing GPU Schedulers

Queue and dispatch [Kato ATC11] [Kato RTSS11]



Bandwidth-aware non-preemptive device (BAND) Scheduler



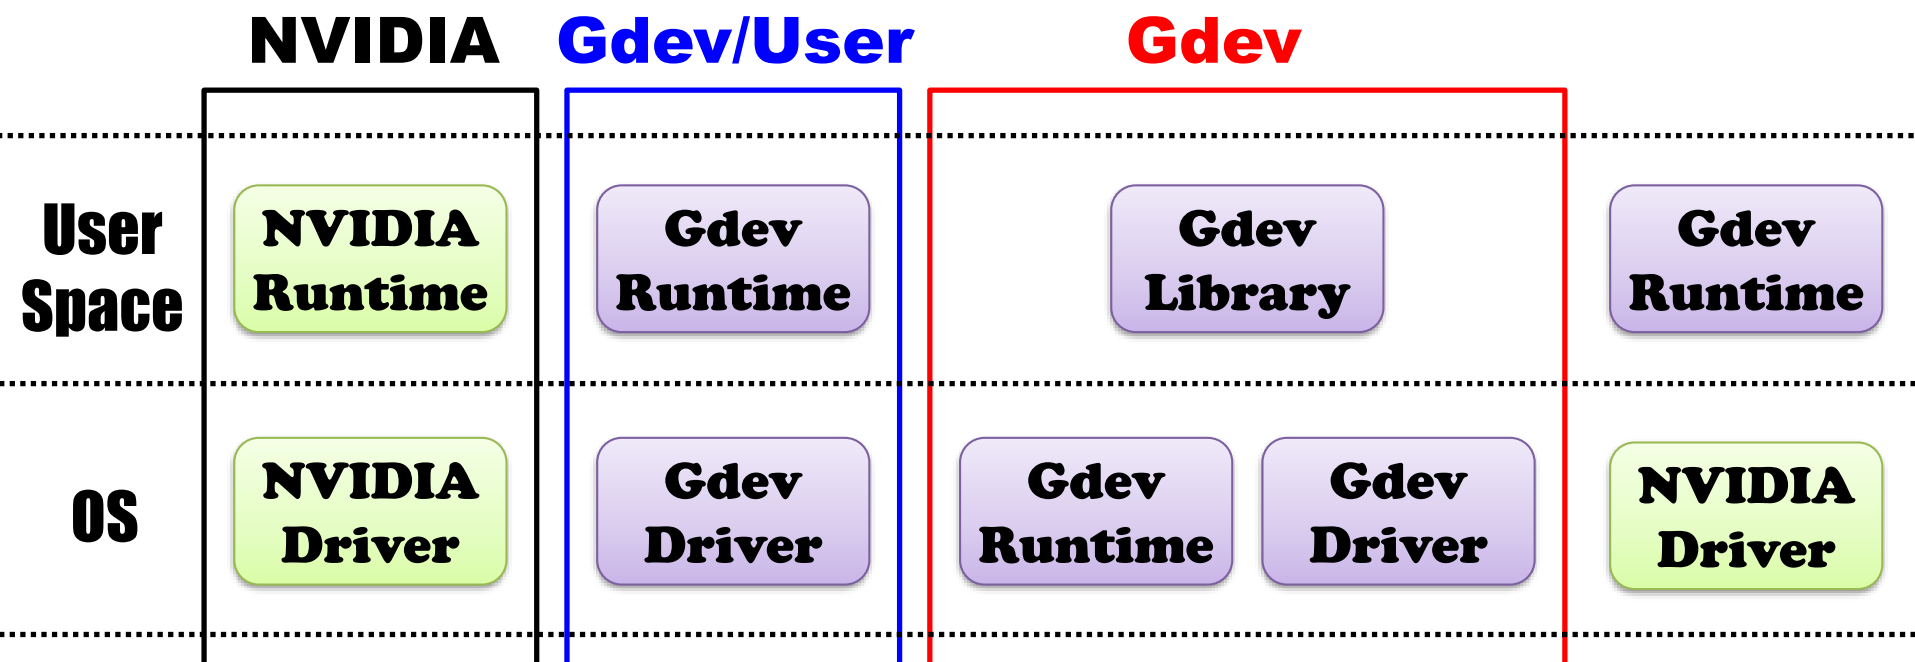
Outline

- **Motivation**
- **Approach**
- **GPU Resource Management**
- **Evaluation**
- **Conclusion**


Experimental Setup

- Linux kernel 2.6.39
- NVIDIA GeForce GTX 480
- Intel Core 2 Extreme QX9650
- NVIDIA CUDA Compiler 4.0 and GCC 4.4.6
- Benchmarks & Applications:
 - Rodinia benchmark [Che *et al*, IISWC'09]
 - eCryptfs encrypted filesystem
 - FAST database search [Kim *et al*, SIGMOD'10]
 - PTask dataflow benchmarks [Rossbach *et al*, SOSPP'11]


Runtime and Driver Choice



Basic Performance


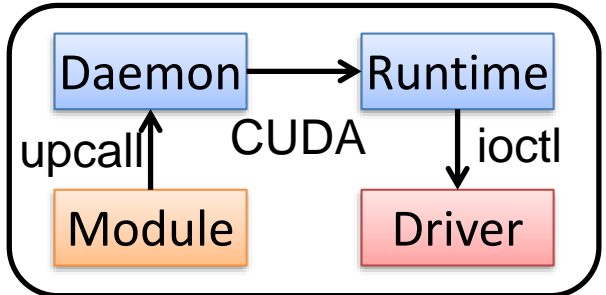


eCryptfs Read&Write Throughput

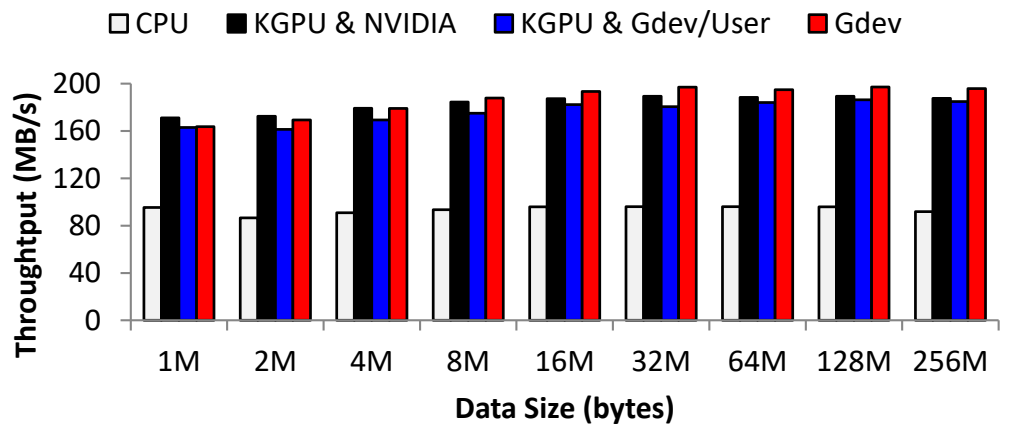


Read throughput


KGPU
Sun et. al.
2012




Write throughput




Impact of Shared Device Memory



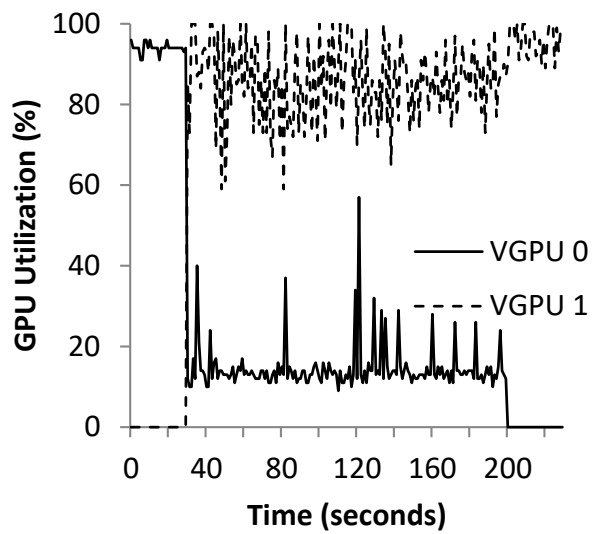
Impact of Data Swapping




Virtual GPU Isolation



No scheduling
(FIFO)



Xen VM Policy
(Credit)



Gdev Policy
(BAND)

Outline

- Motivation
- Approach
- GPU Resource Management
- Evaluation
- **Conclusion**

Concluding Remarks

Gdev is an OS approach to **first-class** GPU resource management.

GPUs can be **used by the OS**.

GPUs can be **protected by the OS**.

GPUs can be **multi-tasked by the OS**.

Compromising basic performance to some extent.

Concluding Remarks


Gdev is **open-source**.

Facilitate systems research.

Visit <http://sys.ertl.jp/gdev/>.

What's up-to-date:

- RAID6 erasure coding acceleration.
- Dynamic power management.
- Zero-copy between I/O devices and GPUs.



Thank You!

Questions?

